

1. OpenAI — GPT-5.5

Release: 23 April 2026

Knowledge Cutoff: Nov 2025

Context Window: ~1.1M tokens

Pricing (API):

- Input: \$5 / 1M tokens
- Output: \$30 / 1M tokens

Availability: Public (ChatGPT, API)

Why / When to use:

Best-in-class reasoning + agentic workflows

- Multi-step planning, coding agents, long-document analysis
- High context + strong tool use

2. Anthropic — Claude Sonnet 4.6

Release: February 2026

Knowledge Cutoff: Not publicly disclosed

Context Window: ~1M tokens

Pricing (API):

- Input: \$3 / 1M tokens
- Output: \$15 / 1M tokens

Availability: Public (Claude, API)

Why / When to use:

Best for long-context coding + enterprise AI

- Large codebase analysis
- Enterprise document workflows
- High-quality summarization and reasoning

3. Google DeepMind — Gemini 3.1 Pro

Release: 2026

Knowledge Cutoff: Not publicly disclosed

Context Window: ~1M tokens

Pricing (API):

- Input: ~\$1.25 / 1M tokens
- Output: ~\$10 / 1M tokens

Availability: Public (Gemini API, Vertex AI)

Why / When to use:

Best for multimodal reasoning + Google ecosystem

- Image + text + video workflows

- Search-heavy reasoning
- Enterprise AI on Google Cloud

4. xAI — Grok 3

Release: February 2025

Knowledge Cutoff: Frequently refreshed

Context Window: ~128K–256K tokens

Pricing (API):

- Usage-based pricing
- Varies by deployment

Availability: Public (X platform, API)

Why / When to use:

Best for real-time web + conversational AI

- News summarization
- Social media analysis
- Fast conversational responses

5. Meta — Llama 4 Maverick

Release: April 2025

Knowledge Cutoff: Not publicly disclosed

Context Window: ~1M+ tokens

Pricing (API):

- Open-weight model
- Hosting cost depends on provider

Availability: Public (Open-weight)

Why / When to use:

Best for self-hosting + fine-tuning

- On-prem AI deployment
- Research experimentation
- Cost-efficient inference

6. Mistral AI — Magistral Medium

Release: 2026

Knowledge Cutoff: Not publicly disclosed

Context Window: ~128K+ tokens

Pricing (API):

- Enterprise pricing
- Custom deployment pricing

Availability: Limited/Public API

Why / When to use:

Best for European enterprise AI

- Multilingual workflows
- EU compliance-focused deployments
- Efficient enterprise reasoning

7. DeepSeek — DeepSeek-R1

Release: January 2025

Knowledge Cutoff: Not publicly disclosed

Context Window: ~128K tokens

Pricing (API):

- Very low-cost inference
- Cheaper than most frontier models

Availability: Public

Why / When to use:

Best for low-cost reasoning + coding

- Budget AI deployments
- Coding copilots
- Math-heavy workloads

8. Cohere — Command A

Release: March 2025

Knowledge Cutoff: Not publicly disclosed

Context Window: ~256K tokens

Pricing (API):

- Enterprise/API pricing
- Custom enterprise contracts

Availability: Public API

Why / When to use:

Best for RAG + enterprise assistants

- Enterprise search systems
- Retrieval pipelines
- Multilingual business assistants

9. Alibaba Cloud — Qwen3

Release: 2025

Knowledge Cutoff: Not publicly disclosed

Context Window: Up to ~1M tokens

Pricing (API):

- Open-weight + API variants
- Competitive low-cost pricing

Availability: Public

Why / When to use:

Best for multilingual + coding workloads

- Chinese + English AI systems
- Open-source customization
- Software engineering tasks

10. AI21 Labs — Jamba 1.6 Large

Release: 2025

Knowledge Cutoff: Not publicly disclosed

Context Window: ~256K+ tokens

Pricing (API):

- Enterprise/API pricing
- Custom deployment pricing

Availability: Public API

Why / When to use:

Best for long-context enterprise workflows

- Legal/document-heavy pipelines
- Structured output generation
- AI orchestration systems